

The World Wide Web

COMPUTER NETWORKS A.A. 24/25



Leonardo Maccari, DAIS: Ca' Foscari University of Venice,
leonardo.maccari@unive.it

Venice, fall 2024

Sect. 1 WWW



The World Wide Web



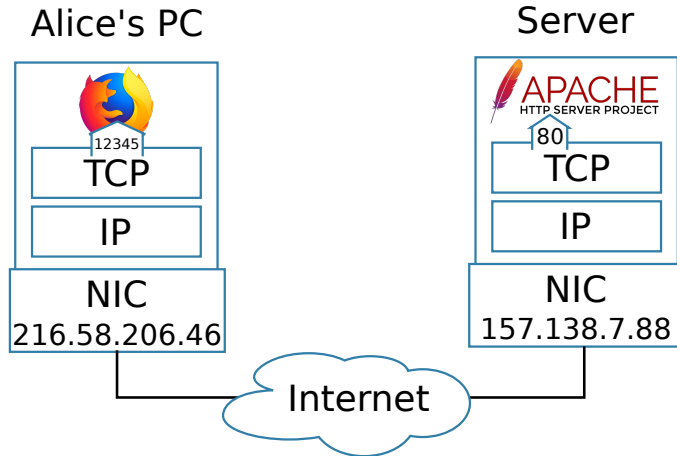
What we know today as the WWW is a mix of various components

- A schema for creating references to locations
- A language used to build pages
- An application layer protocol to retrieve the pages


Server Organization



- Before we go into the details of the protocols, let's see how the application functionally work

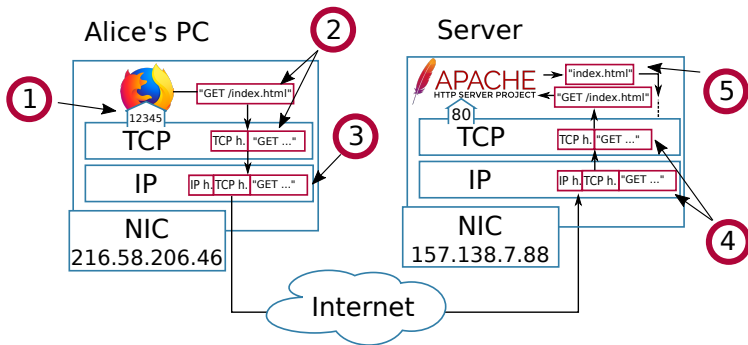


Communication Flow



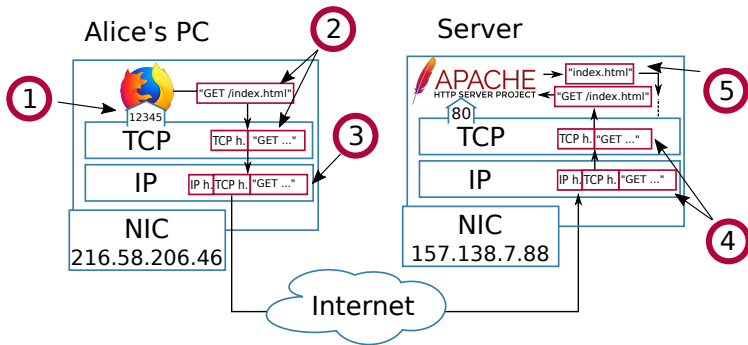
1. The browser is assigned a port from the TCP layer, the web server gets port 80

Communication Flow



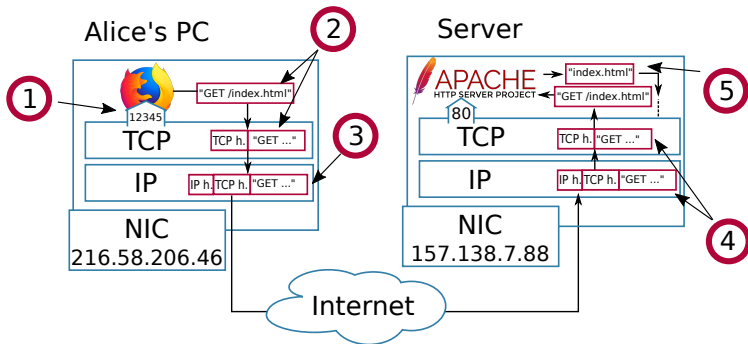
2. The browser issues an HTTP command (more on this later), it gets encapsulated in a TCP segment

Communication Flow




3. This gets encapsulated in an IP segment

Communication Flow



4. The IP and TCP layer decapsulate it

Communication Flow



5. The command is received by the web server, that elaborates and answers

WWW

↳ 1.1 URI

Universal Resource Identifiers: URI



- An URI is a character string that unambiguously identifies a resource on the world wide web
- An URL is an URI that provides you with the location to retrieve a resource, contrarily to an URN that is an URI that simply identifies its name, and not its location
- A (simplified) format of an URI/URL is the following:

URI = scheme ":" "://" authority path ["?" query] ["#" fragment]

- let's see it piece by piece

- A scheme can be seen as a selector, indicating the meaning of the fields after it.
- In practice, the scheme often identifies the application-layer protocol that must be used by the client to retrieve the document, but it is not always the case
- Some valid schemes are: `http`, `https`, `ftp`, `mailto`

URI: authority



- This includes the domain name or the IP address of the server where the document can be retrieved using the protocol specified via the scheme.
- This name can be preceded by some information about the user (e.g. a user name and a password, even if passwords in URIs are deprecated) who is requesting the information.
- This is included before the @ character
- The host name can be followed by the colon character (:) and a port number.
- The port should be included only if the scheme does not implicitly sets one.

URI: path



- The third part of the URI is the path to the document.
- This path is structured as filenames on a Unix host (but it does not imply that the files are indeed stored this way on the server).
- If the path is not specified, the server will return a default document.



URI: query and fragment



- The last two optional parts of the URI are used to provide a query parameter, that will be parsed by the code on the receiving server (?parameter=value)
- Several parameters are separated using the ampersand &
- and to indicate a specific fragment in the document #SubsectionTitle to point to

URI examples



`http://tools.ietf.org/html/rfc3986.html`

schema: `http`

authority: `tools.ietf.org`

path: `/html/rfc3986.html`

URI examples



`mailto:infobot@example.com?subject=current-issue`

schema: `mailto`

authority: `infobot` (user) in the `example.com` domain

path: `none`

query: `subject=current-issue`

This will be opened by an a MUA that will start editing an e-mail for `infobot@example.com` with subject *current issue*.

URI examples



`https://en.wikipedia.org/wiki/Italy#Name`

schema: `http`

authority: `en.wikipedia.org`

path: `/wiki/Italy`

query: `none`

fragment: `#Name`. The page will be displayed starting from this section

URI examples



`ftp://cnn.example.com&story=breaking_news@10.0.0.1/top_story.htm`

And what about this?



Sect. 2 HTML

- Now that we know how to refer to objects, we need a way to add links to documents
- To do this we enrich the text inside documents with attributes, named tags
- The way this is done is with a Markup language, The Hypertext Markup Language (HTML)

```
1 <tag>Some text to be displayed</tag>
2
3 <tag attribute1="value1" attribute2="value2">some text to be displayed</tag>
```

- HTML was first standardized by IETF, then its following versions were taken over by the W3C consortium
- Right now the latest stable version is HTML 5.0

An HTML page



- An HTML document is made of a header and a body, included in the respective tags
- Inside the page the text can be formatted using tags, including the presence of external objects
- These external objects will be referenced by their URL and will be loaded separately by the Browser, when it renders the page

Example Page



In the previous example some common tags are shown, like:

IMG: This provides the URL for a remote image. The browser will load the image and compose it with the rest of the page. Images can be stored anywhere on the Internet

HR: an horizontal rule

UL, LI: Unordered list and a list item

A: a link, where **HREF** is the URI

HTML

↳ 2.1 Dynamic Pages

Dynamic Pages in HTML



- An HTML page is just a file, so if you want to have a page with dynamic contents, there are two ways.
- The first is to add server-side code, that makes the page change when you visit it, based on what you have done before, or based on your query.
- This means that the URL does not point to an HTML file, but to a file in some language that gets executed on the server, and will return an HTML file to the browser.
- This is referred to as Common Gateway Interface: CGI

PHP Example



PHP is a widely used server-side programming language, that is embedded inside the HTML page itself, or it is directly referenced by the URL, when you see it points to a .php file.

```
1 <html>
2   <body>
3     <h1>Example PHP page</h1>
4     <?php echo "This is the output of PHP code"; ?>
5   </body>
6 </html>
```

This code will produce something like:

Example PHP page

This is the output of PHP code

PHP query variables



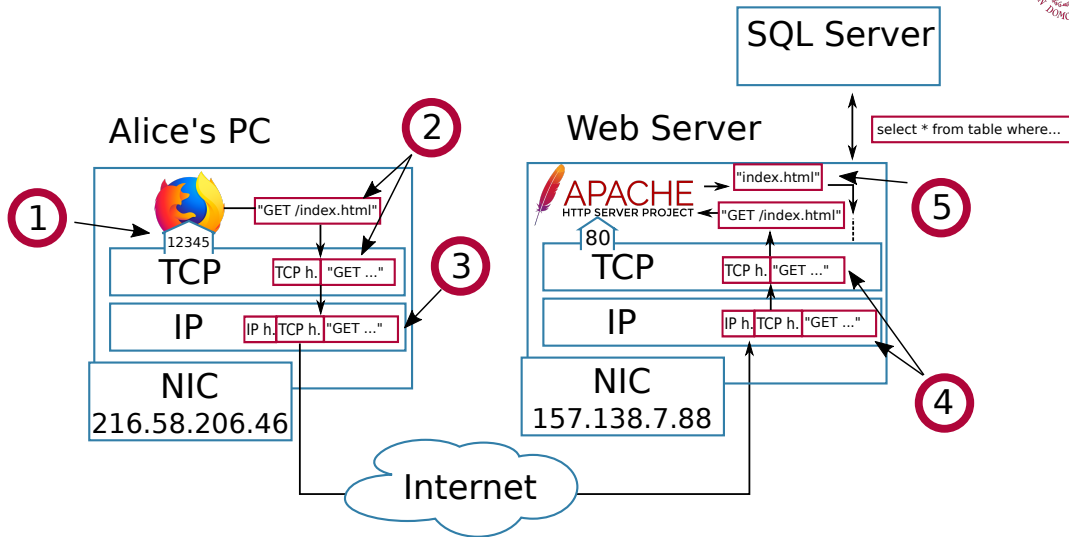
- Server-side code can use the variables you pass as queries in the URL
- For instance using the following code, if the user browses:

`http://example.com/?name=Hannes`

```
1 <html>
2 <body>
3 <?php
4     echo "Hello " . htmlspecialchars($_GET["name"]) . "!";
5 ?> </body>
6 </html>
```

This code will produce something like:

Hello Hannes!



- With server-side programs the user should not access the source code.
- The code is ran by the server, the HTML is generated and delivered to the browser
- The other way of making dynamic pages is with client-side programming
- This means that the code is embedded in the HTML page, sent to the browser and the browser runs it.
- The most popular client-side language is Javascript

Javascript Example



```
1 <html>
2 <body>
3   <p>Hello World!</p>
4   <script>
5     alert("Hello , world! in JS");
6   </script>
7 </body>
8 </html>
```

Try it yourself. Save the code in an HTML file and open it with your browser.

Client Vs Server side code



- The two models are used for similar but different things
- The server-side code is used for instance, to query a database and select the right information. The user should not have access to the database itself, only to the information it contains
- The client-side is used to make dynamic actions, that in some cases may even not require the use of the connection
- You can use both things together



Example



- An e-commerce has a database of thousands of items you can purchase
- When you search for a specific keyword, the back-end will filter all the objects that are relevant for you and send them to the browser
- The browser then can order them based on the user desire (by price, age. . .). This will be done client-side by JS in the user browser, without the need to query the website again.



HTML

↳ 2.2 Style Sheets

Formatting the Page content: CSS



- The Cascade Style Sheet CSS is another standard that allows to graphically adjust the contents of your web page
- It provides decoupling between data and style.
- For instance, you want to write some text in red.
- You can do it inside the HTML page, but then...

CSS Example



Try to save a file like this, and browse it

```
1 <html>
2 <body>
3   <h1>This is a title</h1>
4   <p style="color:#00FF00">This is a simple paragraph</p>
5 </body>
6 </html>
```

This is what you should see



This is a title


This is a simple paragraph

- If you want to modify the style of all paragraphs, you need to modify all the `<p>` tags
- CSS instead allow to specify the style in a different file, that is loaded by the browser with a `<link>` tag when the page is served
- On the left you have the `example.html` file. It includes the `style.css` file, that is on the right.

```
1 <html>
2 <link href="style.css" rel="stylesheet" type="text/css">
3 <body>
4   <h1>This is a title</h1>
5   <p>This is a simple paragraph</p>
6 </body>
7 </html>
```

```
1 h1 {
2   color: black;
3   font-size: 5em;
4 }
5 p {
6   color: green;
7 }
```

The Styled HTML



This is a title

This is a simple paragraph

A melting Pot of *includes*



- Nowadays many sites load their components from a multitude of other sources, such as:
 - They load JS libraries from third parties, to exploit the newest features
 - They load CSS from known sites like Twitter Bootstrap or FontAwesome
 - They load images from external sites and caches
- Try it yourself. Open the Chrome browser, then "more tools"→"Developer tools"→"sources" (and "network") and load some popular web page like `repubblica.it`

Sect. 3 HyperText Transfer Protocol HTTP

- The third component of the world wide web is the HyperText Transfer Protocol (HTTP), a text-based protocol like SMTP.
- The client sends a request and the server returns a response.
- HTTP runs above the bytestream service (TCP) and HTTP servers listen by default on port 80.

- Each HTTP request contains three parts :
 - a method, that indicates the type of request, a URI, and the version of the HTTP protocol used by the client
 - a header, that is used by the client to specify optional parameters for the request. An empty line is used to mark the end of the header
 - an optional MIME document attached to the request

- The response sent by the server also contains three parts :
 - a status line, that indicates whether the request was successful or not
 - a header, that contains additional information about the response. The response header ends with an empty line.
 - a MIME document

The most popular methods are

- GET: used to retrieve a document from a server. It is followed by the path of the URI of the requested document and the version of HTTP used by the client.
- HEAD: allows the retrieval of the header lines for a given URI without retrieving the entire document. It can be used by a client to verify if a document exists, for instance.
- POST: used by a client to send a document to a server. The document is attached to the HTTP request as a MIME document.

GET example



For example, to retrieve the `http://www.w3.org/MarkUp/` URI, a client must:

- open a TCP connection on port 80 with host `www.w3.org`
- send a HTTP request containing the following line:

```
GET /MarkUp/ HTTP/1.0
```

HTTP headers



MIME headers are partly the same as in the SMTP protocol, for instance the ones that express the encoding. Some new ones are:

User-Agent: some info the client tells about the software running on the client itself. Used by servers to provide different versions of the website to different devices¹

Referrer: the previous website that was browsed by this client. Very useful to understand what are the pages that link to your page and send you the more traffic.

Host: the domain the client requests, in case on the server there is more than one domain, with the same internal path (like /index.html)

Server: some info about the software running on the server. Not really used anymore, in order to avoid security scanners.

¹Not a very good idea https://developer.mozilla.org/en-US/docs/Web/HTTP/Browser_detection_using_the_user_agent

- The status line of the HTTP response includes the HTTP version followed by a three digit status code and additional information in English.
- HTTP status codes have a similar structure as the reply codes used by SMTP:
 - 2** Valid response. 200 Ok indicates that the HTTP request was OK.
 - 3** Document is no longer available on the server.
 - 4** the server has detected an error in the HTTP request sent by the client. 400 Bad Request indicates a syntax error in the HTTP request. 404 Not Found indicates that the requested document does not exist on the server.
 - 5** All status codes starting with digit 5 indicate an error on the server.

Example



```
1 GET / HTTP/1.0
2 User-Agent: curl/7.19.4 (universal-apple-darwin10.0) libcurl/7.19.4 OpenSSL/0.9.8l zlib/1.2.3
3 Host: www.ietf.org
```

```
1 HTTP/1.1 200 OK
2 Date: Mon, 15 Mar 2010 13:40:38 GMT
3 Server: Apache/2.2.4 (Linux/SUSE) mod/_ssl/2.2.4 OpenSSL/0.9.8e (truncated)
4 Last-Modified: Tue, 09 Mar 2010 21:26:53 GMT
5 Content-Length: 17019
6 Content-Type: text/html
7 <!DOCTYPE HTML PUBLIC .../HTML>
```

- Initially, one TCP connection was available for only one HTTP request, which was OK for simple pages
- Today we have HTML pages made of hundreds of parts (images, icons, CSS...) and every one requires a different HTTP request.
- So to download a web page with 10 images, we need one for the page and one for each image.
- The Connection: Keep Alive header was added in order to allow to send more than one request in one connection only.
- The client sends the header, and the server answers with the same header specifying how many requests can be used in the rest of the connection

Example



```
1 GET / HTTP/1.1
2 Host: www.kame.net
3 User-Agent: Mozilla/5.0 (Macintosh; U; Intel Mac OS X 10_6_2; en-us)
4 Connection: Keep-Alive
```

```
1 HTTP/1.1 200 OK
2 Date: Fri, 19 Mar 2010 09:23:37 GMT
3 Server: Apache/2.0.63 (FreeBSD) PHP/5.2.12 with Suhosin-Patch
4 Keep-Alive: timeout=15, max=100
5 Connection: Keep-Alive
6 Content-Length: 3462
7 Content-Type: text/html
8 <html>...
9 </html>
```

Stateless operation



- It is important to know that every HTTP request, being or not in the same connection is uncorrelated from the previous ones.
- The HTTP protocol does not maintain a state, or some link from one request to another
- This is different than saying that the backend (the software running on the server that creates and provides the pages to the HTTP server) does not maintain a state. In fact it does.



Typical Example: Authentication



- After you perform an authentication, the website offers you different content
- The backend needs to remember that one HTTP request was made by the same client that made the authentication
- The most common way of doing it is with so-called cookies

- A cookie is some information generated by the server and sent to the client
- The client will have to return the cookie at every request, in order to show that it is the same one that received the cookie initially.
- HTTP does not care about cookies, but offers two headers `Set-Cookie` and `Cookie` that can be used to move cookie back and forward between the server and the client.
- Whenever the client goes back to the same website, it may provide the cookie to be recognized as the same user

- A cookie itself is just a seemingly random number, however,
- . . . entire books and international laws have been written on cookies, but their details mostly pertain to web programming and not to networking topics
- Some terminology and some key concepts follow

Types of Cookies



Cookies can be distinguished on the way they are used:

- Session management: user log-in
- Personalization: client language or geographical location
- Tracking: profiling user activities



A Typical Attack



- If Alice has authenticated to Bob's website and received a cookie, if Eve can access the cookie, it can impersonate Alice
- This is the case in which Eve is able, for instance, to poison the cache of a DNS server, and convince Alice to connect to Eve site instead of Bob's even only for some time
- Alice will connect to Eve site, provide the cookie, and Eve can impersonate Alice with Bob.

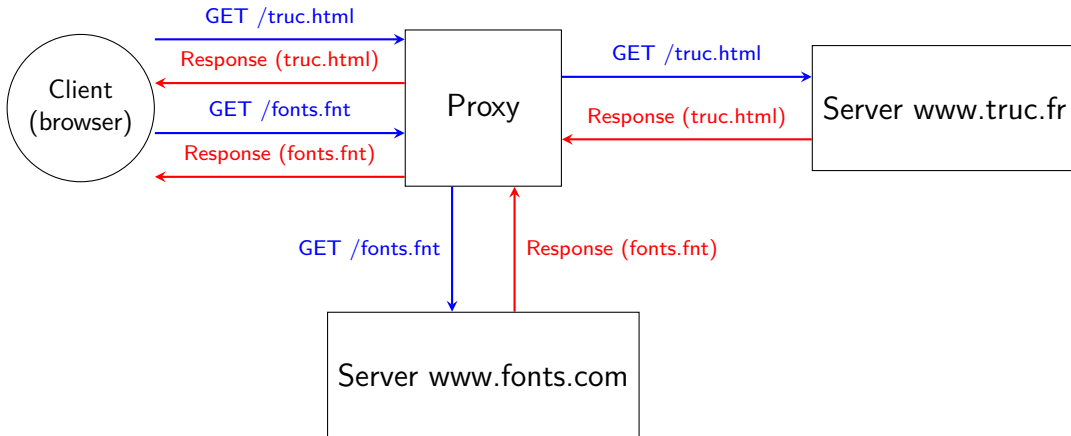
HyperText Transfer Protocol HTTP

↳ 3.1 Proxies and HTTP 2.0


- With time passing, the initial design of HTTP proved to be too slow for modern requirements
- In particular, a sequence of iterative calls to the server is too much depending on the Round-Trip-Time of the connection and so loading a page can take many seconds
- Plus, HTTP server now use dynamic pages and create pages from content included in databases, and the elaboration time it takes to retrieve all components is another source of delay
- We addressed the second issue with proxies, and the first with a totally new version of HTTP

- A proxy server is a cache that saves browsed pages for a certain amount of time
- It resides in the user network
- Browsers need to be configured to use the proxy and so the browser, instead of connecting to the destination HTTP server, will connect to the proxy
- The proxy will then connect to the original HTTP server, retrieve the content and save it for a while
- If the browser navigates the website again the proxy will not start a fresh connection but will simply serve the same page again

Proxy Servers



Proxy Servers



- In the example the web page is made of some HTML content and some fonts, coming from two different servers
- Once the proxy cached the whole page, it will serve it to further requests
- This has two benefits: speeds up the access to the page (no iterative requests) and saving in downlink bandwidth

- The proxy is set-up by the network administrator and can be accessed by the terminals in the local network
- It is a real MiTM so it must be carefully managed



Some HTTP headers are used to know if the page has changed, and needs to be reloaded.

Cache-Control: This is used by the server to notify if the page can not be cached (`no-store`), can be cached up to a certain number of seconds (`max-age=10`), or it can be cached, but follow up requests must still ask to the server if some change happened using the **If-Modified-Since:** client header (`no-cache`)

If-Modified-Since: if client requests a page using this header, it is returned only if the page was modified since a certain time. Else it is not returned by the server, the cache is still valid.

- A reverse proxy is a proxy that is placed on the server premises instead that on the client ones
- The reverse proxy *hides* the real server, so that the client does not know the proxy exists
- The reverse proxy will query the servers and return the pages, exactly as the other kind of server.
- The goal of the reverse proxy is not to save bandiwdth, but to reduce the queries to the databases and the CPU load on the servers.

HyperText Transfer Protocol HTTP

↳ 3.2 HTTP 2.0

- The most recent version of HTTP changed some of the internal ways HTTP works, with the goal of speeding up and changing some of the logic of HTTP
- The key features of HTTP 2.0 are the following ones. . .



A binary protocol



- HTTP 2.0 is not a text protocol anymore
- Now requests and replies are encoded in binary format
- Textual protocols are easy to debug and understand, but slow in practice, text can be largely compressed

- The content of the page is split in *frames*
- This has nothing to deal with the (now obsolete) frames in HTML
- Frames can be sent in parallel data streams from the server to the client, so there is not anymore a strict sequence of requests that the client will make
- This allows the server to send smaller objects before larger ones (and not in the sequence they were asked by the client)

- *head-of-line* blocking is a generic term to indicate a situation in which a large job slows down many small ones
- when the browser loads an HTML page it shows the content as soon as it gets downloaded
- So if the first content to be requested is a big image that takes time to be downloaded, this will block all the other elements and the web page will stop loading
- With HTML 2.0 it's the server to decide which, among all the elements that the client requested, is the one to prioritize, so it can give priority to small textual ones.



- HTTP 1.1 relied on a fully request-reply architecture
- in HTTP 2.0 the server is allowed to *push* content to the client, even if the server did not ask for it
- This allows to save time, when the server knows that some content is required even before the client requests it.

